


**The verifier is not the problem, the score is.  
Same verifier, same 9.87% false-alarm rate, three scores.**



PROJECTED, not measured: replay cannot regenerate text, so repair is modelled. A false alarm can only fire on a step the gate chose to verify, and a good score rarely chooses a correct one.